

# Example Programs for CVODES

## v2.8.2

Radu Serban and Alan C. Hindmarsh  
*Center for Applied Scientific Computing*  
*Lawrence Livermore National Laboratory*

July 30, 2015



UCRL-SM-208115

## **DISCLAIMER**

This document was prepared as an account of work sponsored by an agency of the United States government. Neither the United States government nor Lawrence Livermore National Security, LLC, nor any of their employees makes any warranty, expressed or implied, or assumes any legal liability or responsibility for the accuracy, completeness, or usefulness of any information, apparatus, product, or process disclosed, or represents that its use would not infringe privately owned rights. Reference herein to any specific commercial product, process, or service by trade name, trademark, manufacturer, or otherwise does not necessarily constitute or imply its endorsement, recommendation, or favoring by the United States government or Lawrence Livermore National Security, LLC. The views and opinions of authors expressed herein do not necessarily state or reflect those of the United States government or Lawrence Livermore National Security, LLC, and shall not be used for advertising or product endorsement purposes.

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344.

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Forward sensitivity analysis example problems</b>	<b>6</b>
2.1	A serial nonstiff example: cvsAdvDiff_FSA_non . . . . .	6
2.2	A serial dense example: cvsRoberts_FSA_dns . . . . .	9
2.3	A parallel example with user preconditioner: cvsDiurnal_FSA_kry_p . . . . .	13
<b>3</b>	<b>Adjoint sensitivity analysis example problems</b>	<b>18</b>
3.1	A serial dense example: cvsRoberts_AS Ai_dns . . . . .	18
3.2	A parallel nonstiff example: cvsAdvDiff_AS Ap_non_p . . . . .	21
3.3	A parallel example using CVBBDPRE: cvsAtmDisp_AS Ai_kry_bbd_p . . . . .	24
<b>4</b>	<b>Parallel tests</b>	<b>29</b>
	<b>References</b>	<b>31</b>



# 1 Introduction

This report is intended to serve as a companion document to the User Documentation of CVODES [2]. It provides details, with listings, on the example programs supplied with the CVODES distribution package.

The CVODES distribution contains examples of the following types: serial and parallel examples of Initial Value Problem (IVP) integration, serial and parallel examples of forward sensitivity analysis (FSA), and serial and parallel examples of adjoint sensitivity analysis (ASA). The names of all the examples are given in the following table.

	Serial examples	Parallel examples
IVP	cvsRoberts_dns cvsRoberts_dnsL cvsRoberts_dns_uw cvsRoberts_klu cvsRoberts_sps cvsAdvDiff_bnd cvsAdvDiff_bndL cvsDiurnal_kry cvsDiurnal_kry_bp cvsDirectDemo_ls cvsKrylovDemo_ls cvsKrylovDemo_prec	cvsAdvDiff_non_p cvsDiurnal_kry_p cvsDiurnal_kry_bbd_p
FSA	cvsRoberts_FSA_dns cvsAdvDiff_FSA_non cvsDiurnal_FSA_kry	cvsAdvDiff_FSA_non_p cvsDiurnal_FSA_kry_p
ASA	cvsRoberts_AS Ai_dns cvsAdvDiff_AS Ai_bnd cvsFoodWeb_AS Ai_kry cvsFoodWeb_AS Ap_kry cvsHessian_AS A FSA	cvsAdvDiff_AS Ap_non_p cvsAtmDisp_AS Ai_kry_bbd_p

With the exception of "demo"-type example files, the names of all the examples distributed with SUNDIALS are of the form `[slv][PbName]_[SA]_[ls]_[prec]_[p]`, where

`[slv]` identifies the solver (for CVODES examples this is `cvs`);

`[PbName]` identifies the problem;

`[SA]` identifies sensitivity analysis examples. This field can be one of: `FSA` for forward sensitivity examples, `AS Ai` for adjoint sensitivity examples using an integral-form model output, or `AS Ap` for adjoint sensitivity examples using an pointwise model output;

`[ls]` identifies the linear solver module used (for examples using functional iteration for the nonlinear system solver, `non` specifies that no linear solver was used);

`[prec]` indicates the CVODES preconditioner module used, `bp` for CVBANDPRE or `bbd` for CVBBDPRE (only if applicable, for examples using a Krylov linear solver);

`[p]` indicates an example using the parallel vector module `NVECTOR_PARALLEL`.

The examples are briefly described next. Note that the CVODES distribution includes all of the CVODE C examples (denoted here as examples for IVP integration). More details on these can be found in the CVODE Example Program document [1].

Supplied in the `srcdir/examples/cvodes/serial` directory are the following serial examples (using the `NVECTOR_SERIAL` module):

- `cvsRoberts_dns` solves a chemical kinetics problem consisting of three rate equations. This program solves the problem with the BDF method and Newton iteration, with the CVDENSE linear solver and a user-supplied Jacobian routine. It also uses the rootfinding feature of CVODES.
- `cvsRoberts_dnsL` is the same as `cvsRoberts_dns` but uses the Lapack implementation of CVDENSE.
- `cvsRoberts_dns_uw` is the same as `cvsRoberts_dns` but demonstrates the user-supplied error weight function feature of CVODES.
- `cvsRoberts_klu` is the same as `cvsRoberts_dns` but uses the KLU sparse direct linear solver.
- `cvsRoberts_sps` is the same as `cvsRoberts_dns` but uses the SuperLUMT sparse direct linear solver (with one thread).
- `cvsAdvDiff_bnd` solves the semi-discrete form of an advection-diffusion equation in 2-D. This program solves the problem with the BDF method and Newton iteration, with the CVBAND linear solver and a user-supplied Jacobian routine.
- `cvsAdvDiff_bndL` is the same as `cvsAdvDiff_bnd` but uses the Lapack implementation of CVBAND.
- `cvsDiurnal_kry` solves the semi-discrete form of a two-species diurnal kinetics advection-diffusion PDE system in 2-D. The problem is solved with the BDF/GMRES method (i.e. using the CVSPGMR linear solver) and the block-diagonal part of the Newton matrix as a left preconditioner. A copy of the block-diagonal part of the Jacobian is saved and conditionally reused within the preconditioner setup routine.
- `cvsDiurnal_kry_bp` solves the same problem as `cvsDiurnal_kry`, with the BDF/GMRES method and a banded preconditioner, generated by difference quotients, using the module CVBANDPRE. The problem is solved twice: with preconditioning on the left, then on the right.
- `cvsDirectDemo_ls` is a demonstration program for CVODES with direct linear solvers. Two separate problems are solved using both the Adams and BDF linear multistep methods in combination with functional and Newton iterations. The first problem is the Van der Pol oscillator for which the Newton iteration cases use the following types of Jacobian approximations: (1) dense, user-supplied, (2) dense, difference-quotient approximation, (3) diagonal approximation. The second problem is a linear ODE with a banded lower triangular matrix derived from a 2-D advection PDE. In this case, the Newton iteration cases use the following types of Jacobian approximation: (1) banded, user-supplied, (2) banded, difference-quotient approximation, (3) diagonal approximation.
- `cvsKrylovDemo_ls` solves the same problem as `cvsDiurnal_kry`, with the BDF method, but with three Krylov linear solvers: CVSPGMR, CVSPBCG, and CVSPTFQMR.

- `cvsvKrylovDemo_prec` is a demonstration program with the GMRES linear solver. This program solves a stiff ODE system that arises from a system of partial differential equations. The PDE system is a six-species food web population model, with predator-prey interaction and diffusion on the unit square in two dimensions. The ODE system is solved using Newton iteration and the CVSPGMR linear solver (scaled preconditioned GMRES). The preconditioner matrix used is the product of two matrices: (1) a matrix, only defined implicitly, based on a fixed number of Gauss-Seidel iterations using the diffusion terms only; and (2) a block-diagonal matrix based on the partial derivatives of the interaction terms only, using block-grouping. Four different runs are made for this problem. The product preconditioner is applied on the left and on the right. In each case, both the modified and classical Gram-Schmidt options are tested.
- `cvsvRoberts_FSA_dns` solves a 3-species kinetics problem (from `cvsvRoberts_dns`). CVODES computes both its solution and solution sensitivities with respect to the three reaction rate constants appearing in the model. This program solves the problem with the BDF method, Newton iteration with the CVDENSE linear solver, and a user-supplied Jacobian routine. It also uses the user-supplied error weight function feature of CVODES.
- `cvsvAdvDiff_FSA_non` solves the semi-discrete form of an advection-diffusion equation in 1-D. CVODES computes both its solution and solution sensitivities with respect to the advection and diffusion coefficients. This program solves the problem with the option for nonstiff systems, i.e. Adams method and functional iteration.
- `cvsvDiurnal_FSA_kry` solves the semi-discrete form of a two-species diurnal kinetics advection-diffusion PDE system in 2-D space (from `cvsvDiurnal_kry`). CVODES computes both its solution and solution sensitivities with respect to two parameters affecting the kinetic rate terms. The problem is solved with the BDF/GMRES method (i.e. using the CVSPGMR linear solver) and the block-diagonal part of the Newton matrix as a left preconditioner.
- `cvsvRoberts_ASAdns` solves a 3-species kinetics problem (from `cvsvRoberts_dns`). The adjoint capability of CVODES is used to compute gradients of a functional of the solution with respect to the three reaction rate constants appearing in the model. This program solves both the forward and backward problems with the BDF method, Newton iteration with the CVDENSE linear solver, and user-supplied Jacobian routines.
- `cvsvAdvDiff_ASAdns` solves a semi-discrete 2-D advection-diffusion equation (from `cvsvAdvDiff_bnd`). The adjoint capability of CVODES is used to compute gradients of the average (over both time and space) of the solution with respect to the initial conditions. This program solves both the forward and backward problems with the BDF method, Newton iteration with the CVBAND linear solver, and user-supplied Jacobian routines.
- `cvsvFoodWeb_ASAdns` solves a stiff ODE system that arises from a system of partial differential equations (from `cvsvKrylovDemo_prec`). The PDE system is a six-species food web population model, with predator-prey interaction and diffusion on the unit square in two dimensions.

The adjoint capability of CVODES is used to compute gradients of the average (over both time and space) of the concentration of a selected species with respect to the initial conditions of all six species. Both the forward and backward problems are solved with the BDF/GMRES method (i.e. using the CVSPGMR linear solver) and the block-diagonal part of the Newton matrix as a left preconditioner.

- `cvsFoodWeb_ASAP_kry` solves the same problem as `cvsFoodWeb_ASAi_kry`, but computes gradients of the average over space at the *final time* of the concentration of a selected species with respect to the initial conditions of all six species.
- `cvsHessian_ASA_FSA` is an example of using the *forward-over-adjoint* method for computing 2nd-order derivative information, in the form of Hessian-times-vector products.

Supplied in the `srcdir/examples/cvodes/parallel` directory are the following seven parallel examples (using the `NVECTOR_PARALLEL` module):

- `cvsAdvDiff_non_p` solves the semi-discrete form of a 1-D advection-diffusion equation. This program solves the problem with the option for nonstiff systems, i.e. Adams method and functional iteration.
- `cvsDiurnal_kry_p` is a parallel implementation of `cvsDiurnal_kry`.
- `cvsDiurnal_kry_bbd_p` solves the same problem as `cvsDiurnal_kry_p`, with BDF and the GMRES linear solver, using a block-diagonal matrix with banded blocks as a preconditioner, generated by difference quotients, using the module `CVBBDPRE`.
- `cvsAdvDiff_FSA_non_p` is a parallel version of `cvsAdvDiff_FSA_non`.
- `cvsDiurnal_FSA_kry_p` is a parallel version of `cvsDiurnal_FSA_kry`.
- `cvsAdvDiff_ASAP_non_p` solves a semi-discrete 1-D advection-diffusion equation (from `cvsAdvDiff_non_p`).  
The adjoint capability of CVODES is used to compute gradients of the average over space of the solution at the *final time* with respect to both the initial conditions and the advection and diffusion coefficients in the model. This program solves both the forward and backward problems with the option for nonstiff systems, i.e. Adams method and functional iteration.
- `cvsAtmDisp_AS Ai_kry_bbd_p` solves an adjoint sensitivity problem for an advection-diffusion PDE in 2-D or 3-D using the BDF/GMRES method and the `CVBBDPRE` preconditioner module on both the forward and backward phases.  
The adjoint capability of CVODES is used to compute the gradient of the space-time average of the squared solution norm with respect to problem parameters which parametrize a distributed volume source.

In the following sections, we give detailed descriptions of some (but not all) of the sensitivity analysis examples. We do not discuss the examples for IVP integration; for those, the interested reader should consult the CVODE Examples document [1]. Any CVODE program will work with CVODES with only two modifications: (1) the main program should include



the header file `cvodes.h` instead of `cvode.h`, and (2) the loader command must reference `builddir/lib/lib sundials_cvodes.lib` instead of `builddir/lib/lib sundials_cvode.lib`.

We also give our output files for each of the examples described below, but users should be cautioned that their results may differ slightly from these. Differences in solution values may differ within the tolerances, and differences in cumulative counters, such as numbers of steps or Newton iterations, may differ from one machine environment to another by as much as 10% to 20%.

The final section of this report describes a set of tests done with CVODES in a parallel environment (using `NVECTOR_PARALLEL`) on a modification of the `cvSdiurnal kry_p` example.

In the descriptions below, we make frequent references to the CVODES User Guide [2]. All citations to specific sections (e.g. §4.2) are references to parts of that user guide, unless explicitly stated otherwise.

**Note** The examples in the CVODES distribution were written in such a way as to compile and run for any combination of configuration options during the installation of SUNDIALS (see Appendix ?? in the User Guide). As a consequence, they contain portions of code that will not typically be present in a user program. For example, all example programs make use of the variables `SUNDIALS_EXTENDED_PRECISION` and `SUNDIALS_DOUBLE_PRECISION` to test if the solver libraries were built in extended or double precision, and use the appropriate conversion specifiers in `printf` functions. Similarly, all forward sensitivity examples can be run with or without sensitivity computations enabled and, in the former case, with various combinations of methods and error control strategies. This is achieved in these example through the program arguments.

## 2 Forward sensitivity analysis example problems

For all the CVODES examples, any of three sensitivity method options (`CV_SIMULTANEOUS`, `CV_STAGGERED`, or `CV_STAGGERED1`) can be used, and sensitivities may be included in the error test or not (error control set on `TRUE` or `FALSE`, respectively).

The next three sections describe in detail two serial examples (`cvsAdvDiff_FSA_non` and `cvsRoberts_FSA_dns`), and a parallel one (`cvsDiurnal_FSA_kry_p`). For details on the other examples, the reader is directed to the comments in their source files.

### 2.1 A serial nonstiff example: `cvsAdvDiff_FSA_non`

As a first example of using CVODES for forward sensitivity analysis, we treat the simple advection-diffusion equation for  $u = u(t, x)$

$$\frac{\partial u}{\partial t} = q_1 \frac{\partial^2 u}{\partial x^2} + q_2 \frac{\partial u}{\partial x} \quad (1)$$

for  $0 \leq t \leq 5$ ,  $0 \leq x \leq 2$ , and subject to homogeneous Dirichlet boundary conditions and initial values given by

$$\begin{aligned} u(t, 0) &= 0, \quad u(t, 2) = 0 \\ u(0, x) &= x(2 - x)e^{2x}. \end{aligned} \quad (2)$$

The nominal values of the problem parameters are  $q_1 = 1.0$  and  $q_2 = 0.5$ . A system of `MX` ODEs is obtained by discretizing the  $x$ -axis with `MX+2` grid points and replacing the first and second order spatial derivatives with their central difference approximations. Since the value of  $u$  is constant at the two endpoints, the semi-discrete equations for those points can be eliminated. With  $u_i$  as the approximation to  $u(t, x_i)$ ,  $x_i = i(\Delta x)$ , and  $\Delta x = 2/(\text{MX} + 1)$ , the resulting system of ODEs,  $\dot{u} = f(t, u)$ , can now be written:

$$\dot{u}_i = q_1 \frac{u_{i+1} - 2u_i + u_{i-1}}{(\Delta x)^2} + q_2 \frac{u_{i+1} - u_{i-1}}{2(\Delta x)}. \quad (3)$$

This equation holds for  $i = 1, 2, \dots, \text{MX}$ , with the understanding that  $u_0 = u_{\text{MX}+1} = 0$ .

The sensitivity systems for  $s^1 = \partial u / \partial q_1$  and  $s^2 = \partial u / \partial q_2$  are simply

$$\begin{aligned} \frac{ds_i^1}{dt} &= q_1 \frac{s_{i+1}^1 - 2s_i^1 + s_{i-1}^1}{(\Delta x)^2} + q_2 \frac{s_{i+1}^1 - s_{i-1}^1}{2(\Delta x)} + \frac{u_{i+1} - 2u_i + u_{i-1}}{(\Delta x)^2} \\ s_i^1(0) &= 0.0 \end{aligned} \quad (4)$$

and

$$\begin{aligned} \frac{ds_i^2}{dt} &= q_1 \frac{s_{i+1}^2 - 2s_i^2 + s_{i-1}^2}{(\Delta x)^2} + q_2 \frac{s_{i+1}^2 - s_{i-1}^2}{2(\Delta x)} + \frac{u_{i+1} - u_{i-1}}{2(\Delta x)} \\ s_i^2(0) &= 0.0. \end{aligned} \quad (5)$$

This problem uses the Adams (non-stiff) integration formula and functional iteration. It is unrealistically simple\*, but serves to illustrate use of the forward sensitivity capabilities in CVODES.

---

\*Increasing the number of grid points to better resolve the PDE spatially will lead to a stiffer ODE for which the Adams integration formula will not be suitable.

The `cvAdvDiff_FSA_non.c` file begins by including several header files, including the main CVODES header file, the `sundials_types.h` header file for the definition of the `realtype` type, and the `NVECTOR_SERIAL` header file for the definitions of the serial `N_Vector` type and operations on such vectors. Following that are definitions of problem constants and a data block for communication with the `f` routine. That block includes the problem parameters and the mesh dimension.

The `main` program begins by processing and verifying the program arguments, followed by allocation and initialization of the user-defined data structure. Next, the vector of initial conditions is created (by calling `NVNew_Serial`) and initialized (in the function `SetIC`). The next code block creates and allocates memory for the CVODES object.

If sensitivity calculations were turned on through the command line arguments, the main program continues with setting the scaling parameters `pbar` and the array of flags `plist`. In this example, the scaling factors `pbar` are used both for the finite difference approximation to the right-hand sides of the sensitivity systems (4) and (5) and in calculating the absolute tolerances for the sensitivity variables. The flags in `plist` are set to indicate that sensitivities with respect to both problem parameters are desired. The array of `NS = 2` vectors `uS` for the sensitivity variables is created by calling `NVCloneVectorArray_Serial` and set to contain the initial values ( $s_i^1(0) = 0.0$ ,  $s_i^2(0) = 0.0$ ).

The next three calls set optional inputs for sensitivity calculations: the sensitivity variables are included or excluded from the error test (the boolean variable `err_con` is passed as a command line argument), the control variable `rho` is set to a value `ZERO = 0` to indicate the use of second-order centered directional derivative formulas for the approximations to the sensitivity right-hand sides, and the array of scaling factors `pbar` is passed to CVODES. Memory for sensitivity calculations is allocated by calling `CVodeSensInit1` which also specifies the sensitivity solution method (`sensi_meth` is passed as a command line argument), and the initial conditions for the sensitivity variables. The problem parameters `p` and the arrays `pbar` and `plist` are passed to `CVodeSetSensParam`.

Next, in a loop over the `NOUT` output times, the program calls the integration routine `CVode`. On a successful return, the program prints the maximum norm of the solution  $u$  at the current time and, if sensitivities were also computed, extracts and prints the maximum norms of  $s^1(t)$  and  $s^2(t)$ . The program ends by printing some final integration statistics and freeing all allocated memory.

The `f` function is a straightforward implementation of Eqn. (3). The rest of the source file contains definitions of private functions. The last two, `PrintFinalStats` and `check_flag`, can be used with minor modifications by any CVODES user code to print final CVODES statistics and to check return flags from CVODES interface functions, respectively.


Results generated by `cvAdvDiff_FSA_non` are shown in Fig. 1. The output generated by `cvAdvDiff_FSA_non` when computing sensitivities with the `CV_SIMULTANEOUS` method and full error control (`cvAdvDiff_FSA_non -sensi sim t`) is as follows:

```

----- cvAdvDiff_FSA_non sample output -----
1-D advection-diffusion equation, mesh size = 10
Sensitivity: YES ( SIMULTANEOUS + FULL ERROR CONTROL )

=====
      T      Q      H      NST      Max norm
=====
5.000e-01  4  7.656e-03  115

```



				Sensitivity 2	3.8674e-04
-----					
4.500e+00	3	4.744e-03	680		
				Solution	1.3422e-04
				Sensitivity 1	1.4548e-03
				Sensitivity 2	1.2589e-04
-----					
5.000e+00	1	4.010e-03	757		
				Solution	3.8656e-05
				Sensitivity 1	4.6451e-04
				Sensitivity 2	4.0616e-05
-----					
Final Statistics					
nst	=	757			
nfe	=	1373			
netf	=	1	nsetups	=	0
nni	=	1369	ncfn	=	117
nfSe	=	2746	nfeS	=	5492
netfs	=	0	nsetupsS	=	0
nniS	=	0	ncfnS	=	0

## 2.2 A serial dense example: cvsRoberts\_FSA\_dns

This example is a modification of the chemical kinetics example `cvRoberts_dns` described in [1]. It computes, in addition to the solution of the IVP, sensitivities of the solution with respect to the three reaction rates involved in the model. The ODEs are written as:

$$\begin{aligned}
 \dot{y}_1 &= -p_1 y_1 + p_2 y_2 y_3 \\
 \dot{y}_2 &= p_1 y_1 - p_2 y_2 y_3 - p_3 y_2^2 \\
 \dot{y}_3 &= p_3 y_2^2,
 \end{aligned} \tag{6}$$

with initial conditions at  $t_0 = 0$ ,  $y_1 = 1$  and  $y_2 = y_3 = 0$ . The nominal values of the reaction rate constants are  $p_1 = 0.04$ ,  $p_2 = 10^4$  and  $p_3 = 3 \cdot 10^7$ . The sensitivity systems that are solved together with (6) are

$$\begin{aligned}
 \dot{s}_i &= \begin{bmatrix} -p_1 & p_2 y_3 & p_2 y_2 \\ p_1 & -p_2 y_3 - 2p_3 y_2 & -p_2 y_2 \\ 0 & 2p_3 y_2 & 0 \end{bmatrix} s_i + \frac{\partial f}{\partial p_i}, \quad s_i(t_0) = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}, \quad i = 1, 2, 3 \\
 \frac{\partial f}{\partial p_1} &= \begin{bmatrix} -y_1 \\ y_1 \\ 0 \end{bmatrix}, \quad \frac{\partial f}{\partial p_2} = \begin{bmatrix} y_2 y_3 \\ -y_2 y_3 \\ 0 \end{bmatrix}, \quad \frac{\partial f}{\partial p_3} = \begin{bmatrix} 0 \\ -y_2^2 \\ y_2^2 \end{bmatrix}.
 \end{aligned} \tag{7}$$

The main program is described below with emphasis on the sensitivity related components. These explanations, together with those given for the code `cvRoberts_dns` in [1], will also provide the user with a template for instrumenting an existing simulation code to perform forward sensitivity analysis. As will be seen from this example, an existing simulation code can be modified to compute sensitivity variables (in addition to state variables) by only inserting a few CVODES calls into the main program.

First note that no new header files need be included. In addition to the constants already defined in `cvRoberts_dns`, we define the number of model parameters, `NP` ( $= 3$ ), the number of sensitivity parameters, `NS` ( $= 3$ ), and a constant `ZERO`  $= 0.0$ .

As mentioned in §5.1, the user data structure `data` must provide access to the array of model parameters as the only way for CVODES to communicate parameter values to the right-hand side function `f`. In the `cvsRoberts_FSA_dns` example this is done by defining `data` to be of type `UserData`, i.e. a pointer to a structure which contains an array of `NP` `realtype` values.

Four user-supplied functions are defined. The function `f`, passed to `CVodeInit`, computes the right-hand side of the ODE (6), while `Jac` computes the dense Jacobian of the problem and is attached to the dense linear solver module CVDENSE through a call to `CVDlsSetDenseJacFn`. The function `fS` computes the right-hand side of each sensitivity system (7) for one parameter at a time and is therefore of type `SensRhs1`. Finally, the function `ewt` computes the error weights for the WRMS norm estimations within CVODES.

The program prologue ends by defining six private helper functions. The first two, `ProcessArgs` and `WrongArgs` (which would not be present in a typical user code), parse and verify the command line arguments to `cvsRoberts_FSA_dns`, respectively. After each successful return from the main CVODES integrator, the functions `PrintOutput` and `PrintOutputS` print the state and sensitivity variables, respectively. The function `PrintFinalStats` is called after completion of the integration to print solver statistics. The function `check_flag` is used to check the return flag from any of the CVODES interface functions called by `main`.

The `main` program begins with definitions and type declarations. Among these, it defines the vector `pbar` of `NS` scaling factors for the model parameters `p`, and the array `yS` of vectors (of type `N_Vector`) which will contain the initial conditions and solutions for the sensitivity variables. It also declares the variable `data` of type `UserData` which will contain the user-defined data structure to be passed to CVODES and used in the evaluation of the ODE right-hand sides.

The first code block in `main` deals with reading and interpreting the command line arguments. `cvsRoberts_FSA_dns` can be run with or without sensitivity computations turned on and with different selections for the sensitivity method and error control strategy.

The user's data structure is then allocated and its field `p` is set to contain the values of the three problem parameters. The next block of code is identical to that in `cvRoberts_dns.c` (see [1]) and involves allocation and initialization of the state variables, and creation and initialization of `cvode_mem`, the CVODES solver memory. It specifies that a user-provided function (`ewt`) is to be used for computing the error weights. It also attaches CVDENSE, with a non-NULL Jacobian function, as the linear solver to be used in the Newton nonlinear iteration.

If sensitivity analysis is enabled (through the command line arguments), the main program will then set the scaling parameters `pbar` ( $pbar_i = p_i$ , which can typically be used for nonzero model parameters). Next, the program allocates memory for `yS`, by calling the `NVECTOR_SERIAL` function `N_VCloneVectorArray_Serial`, and initializes all sensitivity variables to 0.0.

The call to `CVodeSensInit1` specifies the sensitivity solution method through the argument `sensi_meth` (read from the command line arguments) as one of `CV_SIMULTANEOUS`, `CV_STAGGERED`, or `CV_STAGGERED1`. It also specifies the user-defined routine, `fS`, for evaluation of the right-hand sides of sensitivity equations.

The next three calls specify optional inputs for forward sensitivity analysis: specifying that sensitivity tolerances are to be based on `pbar`, the error control strategy (read from the

command line arguments), and the information on the model parameters. In this example, only `pbar` is needed for the estimation of absolute sensitivity variable tolerances; neither `p` nor `plist` is required since the sensitivity right-hand sides are computed in the user function `fS`. As a consequence, we pass `NULL` for the corresponding arguments in `CVodeSetSensParams`.

Note that this example uses the default estimates for the relative and absolute tolerances `rtolS` and `atolS` for sensitivity variables, based on the tolerances for state variables and the scaling parameters `pbar` (see §2.6 for details).

Next, in a loop over the `NOUT` output times, the program calls the integration routine `CVode` which, if sensitivity analysis was initialized through the call to `CVodeSensInit1`, computes both state and sensitivity variables. However, `CVode` returns only the state solution at `tout` in the vector `y`. The program tests the return from `CVode` for a value other than `CV_SUCCESS` and prints the state variables. Sensitivity variables at `tout` are loaded into `yS` by calling `CVodeGetSens`. The program tests the return from `CVodeGetSens` for a value other than `CV_SUCCESS` and then prints the sensitivity variables.

Finally, the program prints some statistics (function `PrintFinalStats`) and deallocates memory through calls to `N_VDestroy_Serial`, `N_VDestroyVectorArray_Serial`, `CVodeFree`, and `free` for the user data structure.

The user-supplied functions `f` (for the right-hand side of the original ODEs) and `Jac` (for the system Jacobian) are identical to those in `cvRoberts_dns.c`, with the notable exception that model parameters are extracted from the user-defined data structure `data`, which must first be cast to the `UserData` type. Similarly, the user-supplied function `ewt` is identical to that in `cvRoberts_dns_uw.c`. The user-supplied function `fS` computes the sensitivity right-hand side for the `iS`-th sensitivity equation.

Results generated by `cvsRoberts_FSA_dns` are shown in Fig. 2. The following output is generated by `cvsRoberts_FSA_dns` when computing sensitivities with the `CV_SIMULTANEOUS` method and full error control (`cvsRoberts_FSA_dns -sensi sim t`):

cvsRoberts_FSA_dns sample output						
3-species chemical kinetics problem						
Sensitivity: YES ( SIMULTANEOUS + FULL ERROR CONTROL )						
T	Q	H	NST	y1	y2	y3
4.000e-01	3	4.881e-02	115			
		Solution		9.8517e-01	3.3864e-05	1.4794e-02
		Sensitivity 1		-3.5595e-01	3.9025e-04	3.5556e-01
		Sensitivity 2		9.5431e-08	-2.1309e-10	-9.5218e-08
		Sensitivity 3		-1.5833e-11	-5.2900e-13	1.6362e-11
4.000e+00	5	2.363e-01	138			
		Solution		9.0552e-01	2.2405e-05	9.4459e-02
		Sensitivity 1		-1.8761e+00	1.7922e-04	1.8759e+00
		Sensitivity 2		2.9614e-06	-5.8305e-10	-2.9608e-06
		Sensitivity 3		-4.9334e-10	-2.7626e-13	4.9362e-10
4.000e+01	3	1.485e+00	219			
		Solution		7.1583e-01	9.1856e-06	2.8416e-01
		Sensitivity 1		-4.2475e+00	4.5913e-05	4.2475e+00
		Sensitivity 2		1.3731e-05	-2.3573e-10	-1.3730e-05
		Sensitivity 3		-2.2883e-09	-1.1380e-13	2.2884e-09




Figure 2: Results for the `cvsRoberts_FSA_dns` example problem: time evolution of  $y_1$  and its sensitivities wit-0.247294(t)0.311076(-0.247294(t)0.311076(-621. 5616.62 l 2-621. 5616.661.72 16.62 l 289561 11



				Sensitivity 3	-1.7111e-11	-6.8513e-17	1.7111e-11
-----							
4.000e+07	4	1.776e+06	753				
		Solution		5.2039e-05	2.0817e-10	9.9995e-01	
		Sensitivity 1		-2.5991e-03	-5.1931e-09	2.5991e-03	
		Sensitivity 2		1.0396e-08	2.0772e-14	-1.0397e-08	
		Sensitivity 3		-1.7330e-12	-6.9328e-18	1.7330e-12	
-----							
4.000e+08	4	2.766e+07	802				
		Solution		5.2106e-06	2.0842e-11	9.9999e-01	
		Sensitivity 1		-2.6063e-04	-5.2149e-10	2.6063e-04	
		Sensitivity 2		1.0425e-09	2.0859e-15	-1.0425e-09	
		Sensitivity 3		-1.7366e-13	-6.9467e-19	1.7367e-13	
-----							
4.000e+09	2	4.183e+08	836				
		Solution		5.1881e-07	2.0752e-12	1.0000e-00	
		Sensitivity 1		-2.5907e-05	-5.1717e-11	2.5907e-05	
		Sensitivity 2		1.0363e-10	2.0687e-16	-1.0363e-10	
		Sensitivity 3		-1.7293e-14	-6.9174e-20	1.7293e-14	
-----							
4.000e+10	2	3.799e+09	859				
		Solution		6.5181e-08	2.6072e-13	1.0000e-00	
		Sensitivity 1		-2.4884e-06	-3.3032e-12	2.4884e-06	
		Sensitivity 2		9.9534e-12	1.3213e-17	-9.9534e-12	
		Sensitivity 3		-2.1727e-15	-8.6908e-21	2.1727e-15	
-----							
Final Statistics							
nst	=	859					
nfe	=	1222					
netf	=	29	nsetups	=	142		
nni	=	1218	ncfn	=	4		
nfSe	=	3666	nfeS	=	0		
netfs	=	0	nsetupsS	=	0		
nniS	=	0	ncfnS	=	0		
nje	=	24	nfeLS	=	0		

### 2.3 A parallel example with user preconditioner: cvsDiurnal\_FSA\_kry\_p

As an example of using the forward sensitivity capabilities in CVODES with the Krylov linear solver CVSPGMR and the NVECTOR\_PARALLEL module, we describe a test problem (derived from `cvDiurnal_kry_p`) that solves the semi-discrete form of a two-species diurnal kinetics advection-diffusion PDE system in 2-D space, for which we also compute solution sensitivities with respect to problem parameters ( $q_1$  and  $q_2$ ) that appear in the kinetic rate terms.

The PDE system is

$$\frac{\partial c^i}{\partial t} = K_h \frac{\partial^2 c^i}{\partial x^2} + V \frac{\partial c^i}{\partial x} + \frac{\partial}{\partial y} K_v(y) \frac{\partial c^i}{\partial y} + R^i(c^1, c^2, t) \quad (i = 1, 2), \quad (8)$$

where the superscripts  $i$  are used to distinguish the two chemical species, and where the

reaction terms are given by

$$\begin{aligned} R^1(c^1, c^2, t) &= -q_1 c^1 c^3 - q_2 c^1 c^2 + 2q_3(t) c^3 + q_4(t) c^2, \\ R^2(c^1, c^2, t) &= q_1 c^1 c^3 - q_2 c^1 c^2 - q_4(t) c^2. \end{aligned} \quad (9)$$

The spatial domain is  $0 \leq x \leq 20$ ,  $30 \leq y \leq 50$  (in *km*). The various constants and parameters are:  $K_h = 4.0 \cdot 10^{-6}$ ,  $V = 10^{-3}$ ,  $K_v = 10^{-8} \exp(y/5)$ ,  $q_1 = 1.63 \cdot 10^{-16}$ ,  $q_2 = 4.66 \cdot 10^{-16}$ ,  $c^3 = 3.7 \cdot 10^{16}$ , and the diurnal rate constants are defined as:

$$q_i(t) = \begin{cases} \exp[-a_i / \sin \omega t], & \text{for } \sin \omega t > 0 \\ 0, & \text{for } \sin \omega t \leq 0 \end{cases} \quad (i = 3, 4),$$

where  $\omega = \pi/43200$ ,  $a_3 = 22.62$ ,  $a_4 = 7.601$ . The time interval of integration is  $[0, 86400]$ , representing 24 hours measured in seconds.

Homogeneous Neumann boundary conditions are imposed on each boundary, and the initial conditions are

$$\begin{aligned} c^1(x, y, 0) &= 10^6 \alpha(x) \beta(y), \quad c^2(x, y, 0) = 10^{12} \alpha(x) \beta(y), \\ \alpha(x) &= 1 - (0.1x - 1)^2 + (0.1x - 1)^4 / 2, \\ \beta(y) &= 1 - (0.1y - 4)^2 + (0.1y - 4)^4 / 2. \end{aligned} \quad (10)$$

We discretize the PDE system with central differencing, to obtain an ODE system  $\dot{u} = f(t, u)$  representing (8). In this case, the discrete solution vector is distributed across many processes. Specifically, we may think of the processes as being laid out in a rectangle, and each process being assigned a subgrid of size  $\text{MXSUB} \times \text{MYSUB}$  of the  $x - y$  grid. If there are  $\text{NPEX}$  processes in the  $x$  direction and  $\text{NPEY}$  processes in the  $y$  direction, then the overall grid size is  $\text{MX} \times \text{MY}$  with  $\text{MX} = \text{NPEX} \times \text{MXSUB}$  and  $\text{MY} = \text{NPEY} \times \text{MYSUB}$ , and the size of the ODE system is  $2 \cdot \text{MX} \cdot \text{MY}$ .

To compute  $f$  in this setting, the processes pass and receive information as follows. The solution components for the bottom row of grid points assigned to the current process are passed to the process below it, and the solution for the top row of grid points is received from the process below the current process. The solution for the top row of grid points for the current process is sent to the process above the current process, while the solution for the bottom row of grid points is received from that process by the current process. Similarly, the solution for the first column of grid points is sent from the current process to the process to its left, and the last column of grid points is received from that process by the current process. The communication for the solution at the right edge of the process is similar. If this is the last process in a particular direction, then message passing and receiving are bypassed for that direction.

The overall structure of `main` is very similar to that of the code `cvsRoberts_FSA_dns` described above, with differences arising from the use of the parallel `NVECTOR` module, `NVECTOR_PARALLEL`. On the other hand, the user-supplied routines in `cvsDiurnal_FSA_kry_p`, `f` for the right-hand side of the original system, `Precond` for the preconditioner setup, and `PSolve` for the preconditioner solve, are identical to those defined in the example program `cvDiurnal_kry_p` described in [1]. The only difference is in the routine `fcalc`, which operates on local data only and contains the actual calculation of  $f(t, u)$ , where the problem parameters are first extracted from the user data structure `data`. The program `cvsDiurnal_FSA_kry_p` defines no additional user-supplied routines, as it uses the `CVODES` internal difference quotient routines to compute the sensitivity equation right-hand sides.




Figure 3: Results for the `cvsDiurnal_FSA_kry_p` example problem: time evolution of  $c_1$  and  $c_2$  at the bottom-left and top-right corners (left) and of their sensitivities with respect to  $q_1$ .

Sample results generated by `cvsDiurnal_FSA_kry_p` are shown in Fig. 3. These results were generated on a  $(2 \cdot 40) \times (2 \cdot 40)$  spatial grid. The following output is generated by `cvsDiurnal_FSA_kry_p` when computing sensitivities with the `CV_SIMULTANEOUS` method and full error control (`mpirun -np 4 cvsDiurnal_FSA_kry_p -sensi sim t`):

cvsDiurnal.FSA.kry.p sample output						
2-species diurnal advection-diffusion problem						
Sensitivity: YES ( SIMULTANEOUS + FULL ERROR CONTROL )						
T	Q	H	NST		Bottom left	Top right
7.200e+03	3	3.544e+01	415			
				Solution	1.0468e+04	1.1185e+04
					2.5267e+11	2.6998e+11
				Sensitivity 1	-6.4201e+19	-6.8598e+19
					7.1177e+19	7.6556e+19
				Sensitivity 2	-4.3853e+14	-5.0065e+14
					-2.4407e+18	-2.7842e+18
1.440e+04	3	4.903e+01	640			
				Solution	6.6590e+06	7.3008e+06
					2.5819e+11	2.8329e+11
				Sensitivity 1	-4.0848e+22	-4.4785e+22
					5.9550e+22	6.7173e+22
				Sensitivity 2	-4.5235e+17	-5.4318e+17
					-6.5419e+21	-7.8316e+21
2.160e+04	2	2.948e+01	1198			

				Solution	2.6650e+07	2.9308e+07
					2.9928e+11	3.3134e+11
				Sensitivity 1	-1.6346e+23	-1.7976e+23
					3.8203e+23	4.4991e+23
				Sensitivity 2	-7.6601e+18	-9.4433e+18
					-7.6459e+22	-9.4502e+22
2.880e+04	3	4.499e+01	1542	Solution	8.7021e+06	9.6501e+06
					3.3804e+11	3.7510e+11
				Sensitivity 1	-5.3375e+22	-5.9188e+22
					5.4487e+23	6.7430e+23
				Sensitivity 2	-4.8855e+18	-6.1040e+18
					-1.7194e+23	-2.1518e+23
3.600e+04	3	1.107e+01	1675	Solution	1.4040e+04	1.5609e+04
					3.3868e+11	3.7652e+11
				Sensitivity 1	-8.6141e+19	-9.5761e+19
					5.2718e+23	6.6030e+23
				Sensitivity 2	-8.4328e+15	-1.0549e+16
					-1.8439e+23	-2.3096e+23
4.320e+04	4	8.016e+01	3146	Solution	-2.5802e-07	-2.3679e-07
					3.3823e+11	3.8035e+11
				Sensitivity 1	-3.4238e+08	-3.8876e+08
					5.2753e+23	6.7448e+23
				Sensitivity 2	5.4404e+03	6.6262e+03
					-1.8454e+23	-2.3595e+23
5.040e+04	4	2.431e+02	3194	Solution	8.0708e-08	7.0502e-08
					3.3582e+11	3.8645e+11
				Sensitivity 1	-4.5550e+10	-3.9697e+10
					5.2067e+23	6.9664e+23
				Sensitivity 2	-5.5727e+08	-5.5845e+08
					-1.8214e+23	-2.4371e+23
5.760e+04	4	2.523e+02	3216	Solution	-1.8957e-11	-1.6762e-11
					3.3203e+11	3.9090e+11
				Sensitivity 1	-1.5940e+06	-1.4035e+06
					5.0825e+23	7.1205e+23
				Sensitivity 2	-2.2867e+02	-2.2722e+02
					-1.7780e+23	-2.4910e+23

6.480e+04	4	2.820e+02	3254			
				Solution	-7.8717e-10	-6.7017e-10
					3.3130e+11	3.9634e+11
				-----		
				Sensitivity 1	-3.4373e+07	-2.1671e+07
					5.0442e+23	7.3274e+23
				-----		
				Sensitivity 2	1.3198e+08	1.3575e+08
					-1.7646e+23	-2.5633e+23
				-----		
7.200e+04	4	3.810e+02	3276			
				Solution	-4.4006e-09	-3.7463e-09
					3.3297e+11	4.0389e+11
				-----		
				Sensitivity 1	4.5106e+09	3.8200e+09
					5.0783e+23	7.6382e+23
				-----		
				Sensitivity 2	2.0264e+05	2.0791e+05
					-1.7765e+23	-2.6721e+23
				-----		
7.920e+04	5	6.406e+02	3291			
				Solution	-1.8328e-11	-1.5666e-11
					3.3344e+11	4.1203e+11
				-----		
				Sensitivity 1	1.3650e+07	1.1729e+07
					5.0730e+23	7.9960e+23
				-----		
				Sensitivity 2	6.1764e+01	6.8476e+01
					-1.7747e+23	-2.7972e+23
				-----		
8.640e+04	5	6.406e+02	3302			
				Solution	-2.0206e-13	-1.7557e-13
					3.3518e+11	4.1625e+11
				-----		
				Sensitivity 1	1.1323e+06	9.7319e+05
					5.1171e+23	8.2142e+23
				-----		
				Sensitivity 2	7.6632e+00	8.2818e+00
					-1.7901e+23	-2.8736e+23
				-----		
Final Statistics						
nst	=	3302				
nfe	=	4387				
netf	=	165	nsetups	=	508	
nni	=	4383	ncfn	=	7	
nfSe	=	8774	nfeS	=	17548	
netfs	=	0	nsetupsS	=	0	
nniS	=	0	ncfnS	=	0	

### 3 Adjoint sensitivity analysis example problems

The next three sections describe in detail a serial example (`cvsRoberts_ASAi_dns`) and two parallel examples (`cvsAdvDiff_ASAP_non_p` and `cvsAtmDisp_ASAi_kry_bbd_p`) that perform adjoint sensitivity analysis. For details on the other examples, the reader is directed to the comments in their source files.

#### 3.1 A serial dense example: `cvsRoberts_ASAi_dns`

As a first example of using CVODES for adjoint sensitivity analysis, we examine the chemical kinetics problem (from `cvsRoberts_FSA_dns`)

$$\begin{aligned}\dot{y}_1 &= -p_1 y_1 + p_2 y_2 y_3 \\ \dot{y}_2 &= p_1 y_1 - p_2 y_2 y_3 - p_3 y_2^2 \\ \dot{y}_3 &= p_3 y_2^2 \\ y(t_0) &= y_0,\end{aligned}\tag{11}$$

for which we want to compute the gradient with respect to  $p$  of

$$G(p) = \int_{t_0}^T y_3 dt,\tag{12}$$

without having to compute the solution sensitivities  $dy/dp$ . Following the derivation in §2.7, and taking into account the fact that the initial values of (11) do not depend on the parameters  $p$ , by (2.20) this gradient is simply

$$\frac{dG}{dp} = \int_{t_0}^T (g_p + \lambda^T f_p) dt,\tag{13}$$

where  $g(t, y, p) = y_3$ ,  $f$  is the vector-valued function defining the right-hand side of (11), and  $\lambda$  is the solution of the adjoint problem (2.19),

$$\begin{aligned}\dot{\lambda} &= -(f_y)^T \lambda - (g_y)^T \\ \lambda(T) &= 0.\end{aligned}\tag{14}$$

In order to avoid saving intermediate  $\lambda$  values just for the evaluation of the integral in (13), we extend the backward problem with the following  $N_p$  quadrature equations

$$\begin{aligned}\dot{\xi} &= g_p^T + f_p^T \lambda \\ \xi(T) &= 0,\end{aligned}\tag{15}$$

which yield  $\xi(t_0) = -\int_{t_0}^T (g_p^T + f_p^T \lambda) dt$  and thus  $dG/dp = -\xi^T(t_0)$ . Similarly, the value of  $G$  in (12) can be obtained as  $G = -\zeta(t_0)$ , where  $\zeta$  is solution of the following quadrature equation:

$$\begin{aligned}\dot{\zeta} &= g = y_3 \\ \zeta(T) &= 0.\end{aligned}\tag{16}$$

The main program and the user-defined routines are described below, with emphasis on the aspects particular to adjoint sensitivity calculations.

The calling program includes the CVODES header files `cvodes.h` for CVODES definitions and interface function prototypes, the header file `cvodes_dense.h` for the CVDENSE linear solver module, the header file `nvector_serial.h` for the definition of the serial implementation of the NVECTOR module, NVECTOR\_SERIAL, and the file `sundials_math.h` for the definition of the `SUNRabs` macro. This program also includes two user-defined accessor macros, `Ith` and `IJth`, that are useful in writing the problem functions in a form closely matching their mathematical description, i.e. with components numbered from 1 instead of from 0. Following that, the program defines problem-specific constants and a user-defined data structure, which will be used to pass the values of the parameters  $p$  to various user routines. The constant `STEPS` defines the number of integration steps between two consecutive checkpoints. The program prologue ends with the prototypes of four user-supplied functions that are called by CVODES. The first two provide the right-hand side and dense Jacobian for the forward problem, and the last two provide the right-hand side and dense Jacobian for the backward problem.

The `main` function begins with type declarations and continues with the allocation and initialization of the user data structure, which contains the values of the parameters  $p$ . Next, it allocates and initializes `y` with the initial conditions for the forward problem, allocates and initializes `q` for the quadrature used in computing the value  $G$ , and finally sets the scalar relative tolerance `reltolQ` and vector absolute tolerance `abstolQ` for the quadrature variables. No tolerances for the state variables are defined since `cvsRoberts.ASAi_dns` uses its own function to compute the error weights for WRMS norm estimates of state solution vectors.

The call to `CVodeCreate` creates the main integrator memory block for the forward integration and specifies the `CV_BDF` integration method with `CV_NEWTON` iteration. The call to `CVodeInit` initializes the forward integration by specifying the initial conditions. The call to `CVodeWftolerances` specifies a function that computes error weights. The next call specifies the optional user data pointer `data`. The linear solver is selected to be CVDENSE through the call to its initialization routine `CVDense`. The user-provided Jacobian routine `Jac` is specified through a call to `CVDlsSetDenseJacFn`.

The next code block initializes quadrature computations in the forward phase, by allocating CVODES memory for quadrature integration (the call to `CVodeQuadInit` specifies the right-hand side `fQ` of the quadrature equation and the initial values of the quadrature variable), setting the integration tolerances for the quadrature variables, and finally including the quadrature variable in the error test.

Allocation for the memory block of the combined forward-backward problem is accomplished through the call to `CVadjInit` which specifies `STEPS = 150`, the number of steps between two checkpoints, and specifies cubic Hermite interpolation.

The call to `CVodeF` requests the solution of the forward problem to `TOUT`. If successful, at the end of the integration, `CVodeF` will return the number of saved checkpoints in the argument `ncheck` (optionally, a list of the checkpoints can be obtained by calling `CVodeGetAdjCheckPointsInfo` and the checkpoint information printed).

The next segment of code deals with the setup of the backward problem. First, a serial vector `yB` of length `NEQ` is allocated and initialized with the value of  $\lambda (= 0.0)$  at the final time (`TB1 = 4.0E7`). A second serial vector `qB` of dimension `NP` is created and initialized to 0.0. This vector corresponds to the quadrature variables  $\xi$  whose values at  $t_0$  will be the components of the desired gradient of  $\partial G / \partial p$  (after a sign change). Following that, the program sets the relative and absolute tolerances for the backward integration.

The CVODES memory for the backward integration is created and allocated by the calls to

the interface routines `CVodeCreateB` and `CVodeInitB` which specify the `CV_BDF` integration method with `CV_NEWTON` iteration, among other things. The dense linear solver `CVDENSE` is then initialized by calling the `CVDenseB` interface routine and specifying a non-NULL Jacobian routine `JacB` and user data `data`.

The tolerances for the integration of quadrature variables, `reltolB` and `abstolQB`, are specified through `CVodeQuadSStolerancesB`. The call to `CVodeSetQuadErrConB` indicates that  $\xi$  should be included in the error test. Quadrature computation is initialized by calling `CVodeQuadInitB` which specifies the right-hand side of the quadrature equations as `fQB`.

The actual solution of the backward problem is accomplished through two calls to `CVodeB` — one for intermediate output at  $t = 40$ , and one for the final time  $T0 = 0$ . At each point, the backward solution `yB` ( $= \lambda$ ) is obtained with a call to `CVodeGetB` and the forward solution with a call to `CVodeGetAdjY`. The values of the quadrature variables  $\xi$  at time  $T0$  are loaded in `qB` by calling the extraction routine `CVodeGetQuadB`. The negative of `qB` gives the gradient  $\partial G/\partial p$ .

The main program then carries out a second backward problem. It calls to `CVodeReInitB` and `CVodeQuadReInitB` to re-initialize the backward memory block for a new adjoint computation with a different final time ( $TB2 = 50$ ). This is followed by two calls to `CVodeB`, one for intermediate output at  $t = 40$  and one for the final values at  $t = 0$ . Finally, the gradient  $\partial G/\partial p$  of the second function  $G$  is printed.

The main program ends by freeing previously allocated memory by calling `CVodeFree` (for the `CVODES` memory for the forward problem), `CVadjFree` (for the memory allocated for the combined problem), and `N_VFree_Serial` (for the various vectors).

The user-supplied functions `f` and `Jac` for the right-hand side and Jacobian of the forward problem are straightforward expressions of its mathematical formulation (11). The function `ewt` is the same as the one for `cvRoberts_dns_uw.c`. The function `fQ` implements (16), while `fB`, `JacB`, and `fQB` are mere translations of the backward problem (14) and (15).

The output generated by `cvsRoberts_ASAi_dns` is shown below.

----- cvsRoberts\_ASAi\_dns sample output -----

Adjoint Sensitivity Example for Chemical Kinetics

-----

```
ODE: dy1/dt = -p1*y1 + p2*y2*y3
      dy2/dt =  p1*y1 - p2*y2*y3 - p3*(y2)^2
      dy3/dt =  p3*(y2)^2
```

```
Find dG/dp for
  G = int_t0^tB0 g(t,p,y) dt
  g(t,p,y) = y3
```

```
Create and allocate CVODES memory for forward runs
Forward integration ... done ( nst = 830 )
```

```
ncheck = 5
```

```
-----
G:          3.9983e+07
-----
```

```
Create and allocate CVODES memory for backward run
```



```

Backward integration from tB0 = 4.0000e+07

-----
returned t: 4.0000e+01
tout: 4.0000e+01
lambda(t): 3.9967e+07 3.9967e+07 3.9967e+07
y(t): 7.1583e-01 9.1856e-06 2.8416e-01
-----

Done ( nst = 267 )

-----
returned t: 0.0000e+00
lambda(t0): 3.9967e+07 3.9967e+07 3.9967e+07
y(t0): 1.0000e+00 0.0000e+00 0.0000e+00
dG/dp: 7.6836e+05 -3.0688e+00 5.1143e-04
-----

Re-initialize CVODES memory for backward run
Backward integration from tB0 = 5.0000e+01

-----
returned t: 4.0000e+01
tout: 4.0000e+01
lambda(t): 2.8959e-01 1.7624e+00 9.3567e+00
y(t): 7.1583e-01 9.1856e-06 2.8416e-01
-----

Done ( nst = 158 )

-----
returned t: 0.0000e+00
lambda(t0): 8.4191e+00 1.6097e+01 1.6097e+01
y(t0): 1.0000e+00 0.0000e+00 0.0000e+00
dG/dp: 1.7341e+02 -5.0591e-04 8.4323e-08
-----

Free memory

```

### 3.2 A parallel nonstiff example: cvsAdvDiff\_ASAP\_non\_p

As an example of using the CVODES adjoint sensitivity module with the parallel vector module NVECTOR\_PARALLEL, we describe a sample program that solves the following problem: Consider the 1-D advection-diffusion equation

$$\begin{aligned}
\frac{\partial u}{\partial t} &= p_1 \frac{\partial^2 u}{\partial x^2} + p_2 \frac{\partial u}{\partial x} \\
0 &= x_0 \leq x \leq x_1 = 2 \\
0 &= t_0 \leq t \leq t_f = 2.5,
\end{aligned} \tag{17}$$

with boundary conditions  $u(t, x_0) = u(t, x_1) = 0$ ,  $\forall t$ , and initial condition  $u(t_0, x) = u_0(x) = x(2 - x)e^{2x}$ . Also consider the function

$$g(t) = \int_{x_0}^{x_1} u(t, x) dx.$$

We wish to find, through adjoint sensitivity analysis, the gradient of  $g(t_f)$  with respect to  $p = [p_1; p_2]$  and the perturbation in  $g(t_f)$  due to a perturbation  $\delta u_0$  in  $u_0$ .

The approach we take in the program `cvsAdvDiff_ASAP_non.p` is to first derive an adjoint PDE which is then discretized in space and integrated backwards in time to yield the desired sensitivities. A straightforward extension to PDEs of the derivation given in §2.7 gives

$$\frac{dg}{dp}(t_f) = \int_{t_0}^{t_f} dt \int_{x_0}^{x_1} dx \mu \cdot \left[ \frac{\partial^2 u}{\partial x^2}; \frac{\partial u}{\partial x} \right] \quad (18)$$

and

$$\delta g|_{t_f} = \int_{x_0}^{x_1} \mu(t_0, x) \delta u_0(x) dx, \quad (19)$$

where  $\mu$  is the solution of the adjoint PDE

$$\begin{aligned} \frac{\partial \mu}{\partial t} + p_1 \frac{\partial^2 \mu}{\partial x^2} - p_2 \frac{\partial \mu}{\partial x} &= 0 \\ \mu(t_f, x) &= 1 \\ \mu(t, x_0) = \mu(t, x_1) &= 0. \end{aligned} \quad (20)$$

Both the forward problem (17) and the backward problem (20) are discretized on a uniform spatial grid of size  $M_x + 2$  with central differencing and with boundary values eliminated, leaving ODE systems of size  $N = M_x$  each. As always, we deal with the time quadratures in (18) by introducing the additional equations

$$\begin{aligned} \dot{\xi}_1 &= \int_{x_0}^{x_1} dx \mu \frac{\partial^2 u}{\partial x^2}, \quad \xi_1(t_f) = 0, \\ \dot{\xi}_2 &= \int_{x_0}^{x_1} dx \mu \frac{\partial u}{\partial x}, \quad \xi_2(t_f) = 0, \end{aligned} \quad (21)$$

yielding

$$\frac{dg}{dp}(t_f) = -[\xi_1(t_0); \xi_2(t_0)]$$

The space integrals in (19) and (21) are evaluated numerically, on the given spatial mesh, using the trapezoidal rule.

Note that  $\mu(t_0, x^*)$  is nothing but the perturbation in  $g(t_f)$  due to a  $\delta$ -function perturbation  $\delta u_0(x) = \delta(x - x^*)$  in the initial conditions. Therefore,  $\mu(t_0, x)$  completely describes  $\delta g(t_f)$  for any perturbation  $\delta u_0$ .

Both the forward and the backward problems are solved with the option for nonstiff systems, i.e. using the Adams method with functional iteration for the solution of the non-linear systems. The overall structure of the `main` function is very similar to that of the code `cvsRoberts_ASai_dns` discussed previously with differences arising from the use of the parallel `NVECTOR` module. Unlike `cvsRoberts_ASai_dns`, the example `cvsAdvDiff_ASAP_non.p` illustrates computation of the additional quadrature variables by appending NP equations to the adjoint system. This approach can be a better alternative to using special treatment of the quadrature equations when their number is too small for parallel treatment.

Besides the parallelism implemented by `CVODES` at the `NVECTOR` level, this example uses MPI calls to parallelize the calculations of the right-hand side routines `f` and `fB` and of the spatial integrals involved. The forward problem has size `NEQ = MX`, while the backward problem has size `NB = NEQ + NP`, where `NP = 2` is the number of quadrature equations in




Figure 4: Results for the `cvsAdvDiff_ASAP_non_p` example problem. The gradient of  $g(t_f)$  with respect to the initial conditions  $u_0$  is shown superimposed over the values  $u_0$ .

(21). The use of the total number of available processes on two problems of different sizes deserves some comments, as this is typical in adjoint sensitivity analysis.986sOru ofheotal

```

dgdP(tf)
[ 1]: -1.502544e-01
[ 2]: -1.098456e-02

mu(t0)
[ 1]: 2.775571e-04
[ 2]: 5.625276e-04
[ 3]: 8.474242e-04
[ 4]: 1.127515e-03
[ 5]: 1.393257e-03
[ 6]: 1.641212e-03
[ 7]: 1.860490e-03
[ 8]: 2.049401e-03
[ 9]: 2.196614e-03
[10]: 2.302527e-03
[11]: 2.356404e-03
[12]: 2.360902e-03
[13]: 2.306966e-03
[14]: 2.199483e-03
[15]: 2.032115e-03
[16]: 1.811731e-03
[17]: 1.535589e-03
[18]: 1.212084e-03
[19]: 8.426858e-04
[20]: 4.366699e-04

```

### 3.3 A parallel example using CVBBDPRE: cvsAtmDisp\_ASai\_kry\_bbd\_p

As a more elaborate example of a parallel adjoint sensitivity calculation, we describe next the program `cvsAtmDisp_ASai_kry_bbd_p` provided with CVODES. This example models an atmospheric release with an advection-diffusion PDE in 2-D or 3-D and computes the gradient with respect to source parameters of the space-time average of the squared norm of the concentration. Given a known velocity field  $v(t, x)$  and source function  $S$ , the transport equation for the concentration  $c(t, x)$  in a domain  $\Omega$  is given by

$$\begin{aligned}
\frac{\partial c}{\partial t} - k \nabla^2 c + v \cdot \nabla c + S &= 0, \text{ in } (0, T) \times \Omega \\
\frac{\partial c}{\partial n} &= g, \text{ on } (0, T) \times \partial\Omega \\
c &= c_0(x), \text{ in } \Omega \text{ at } t = 0,
\end{aligned} \tag{22}$$

where  $\Omega$  is a box in  $\mathbb{R}^2$  or  $\mathbb{R}^3$  and  $n$  is the normal to the boundary of  $\Omega$ . We assume homogeneous boundary conditions ( $g = 0$ ) and a zero initial concentration everywhere in  $\Omega$  ( $c_0(x) = 0$ ). The wind field has only a nonzero component in the  $x$  direction given by a Poiseuille profile along the direction  $y$ .

Using adjoint sensitivity analysis, the gradient of

$$G(p) = \int_0^T \int_{\Omega} c^2 dx dt$$

where  $x_i$  is the location of the source of intensity  $S(x_i) = p_i$ , and  $\lambda$  is solution of the adjoint PDE

$$\begin{aligned} -\frac{\partial \lambda}{\partial t} - k \nabla^2 \lambda - v \cdot \nabla \lambda &= c(t, x), \text{ in } (T, 0) \times \Omega \\ (k \nabla \lambda + v \lambda) \cdot n &= 0, \text{ on } (0, T) \times \partial \Omega \\ \lambda &= 0, \text{ in } \Omega \text{ at } t = T. \end{aligned} \quad (25)$$

The PDE (22) is semi-discretized in space with central finite differences, with the boundary conditions explicitly taken into account by using layers of ghost cells in every direction. If the direction  $x^i$  of  $\Omega$  is discretized into  $m_i$  intervals, this leads to a system of ODEs of dimension  $N = \prod_1^d (m_i + 1)$ , with  $d = 2$ , or  $d = 3$ . The source term  $S$  is parameterized as a piecewise constant function and yielding  $N$  parameters in the problem. The nominal values of the source parameters correspond to two Gaussian sources.

The source code as supplied runs the 2-D problem. To obtain the 3-D version, add a line `#define USE3D` at the top of `main`.

The adjoint PDE (25) is discretized to a system of ODEs in a similar fashion. The space integrals in (23) and (24) are simply approximated by their Riemann sums, while the time integrals are resolved by appending pure quadrature equations to the systems of ODEs.

We use BDF with the CVSPGMR linear solver and the CVBBDPRE preconditioner for both the forward and the backward integration phases. The value of  $G$  is computed on the forward phase as a quadrature, while the components of the gradient  $dG/dp$  are computed as quadratures during the backward integration phase. All quadrature variables are included in the corresponding error tests.

Communication between processes for the evaluation of the ODE right-hand sides involves passing the solution on the local boundaries (lines in 2-D, surfaces in 3-D) to the 4 (6 in 3-D) neighboring processes. This is implemented in the function `f_comm`, called in `f` and `fB` before evaluation of the local residual components. Since there is no additional communication required for the CVBBDPRE preconditioner, a NULL pointer is passed for `gloc` and `glocB` in the calls to `CVBBDPrecInit` and `CVBBDPrecInitB`, respectively.

For the sake of clarity, the `cvsAtmDisp_ASAi_kry_bbd_p` example does not use the most memory-efficient implementation possible, as the local segment of the solution vectors (`y` on the forward phase and `yB` on the backward phase) and the data received from neighboring processes is loaded into a temporary array `y_ext` which is then used exclusively in computing the local components of the right-hand sides.

Note that if `cvsAtmDisp_ASAi_kry_bbd_p` is given any command line argument, it will generate a series of MATLAB files which can be used to visualize the solution. The results of a 2-D simulation and adjoint sensitivity analysis with `cvsAtmDisp_ASAi_kry_bbd_p` on a  $80 \times 80$  grid and  $2 \times 4 = 8$  processes are shown in Fig. 5. Results in 3-D<sup>†</sup>, on a  $80 \times 80 \times 40$  grid and  $2 \times 4 \times 2 = 16$  processes are shown in Figs. 6 and 7. A sample output generated by `cvsAtmDisp_ASAi_kry_bbd_p` for a 2D calculation is shown below.

```

cvsAtmDisp_ASAi_kry_bbd_p sample output

Parallel Krylov adjoint sensitivity analysis example
2D Advection diffusion PDE with homogeneous Neumann B.C.
Computes gradient of G = int_t_Omega ( c_i^2 ) dt dOmega
with respect to the source values at each grid point.
```

<sup>†</sup>The name of the executable for the 3-D version is `cvsAtmDisp_ASAi_kry_bbd_p3D`.




Figure 5: Results for the `cvsAtmDisp_ASAi_kry_bbd_p` example problem in 2D. The gradient with respect to the source parameters is pictured on the left. On the right, the gradient was color-coded and superimposed over the nominal value of the source parameters.




Figure 6: Results for the `cvsAtmDisp_ASAi_kry_bbd_p` example problem in 3D. Nominal values of the source parameters.




Figure 7: Results for the `cvsAtmDisp_ASAt_kry_bbd_p` example problem in 3D. Two isosurfaces of the gradient with respect to the source parameters. They correspond to values of 0.25 (green) and 0.4 (blue).

```

Domain:
  0.000000 < x < 20.000000    mx = 20    npe_x = 2
  0.000000 < y < 20.000000    my = 40    npe_y = 2

Begin forward integration... done.    G = 2.522277e+03

Final Statistics..

lenrw  =   8746      leniw =   212
llrw   =   8656      lliw  =    80
nst    =    99
nfe    =   106      nfel   =   122
nni    =   102      nli    =   122
nsetups =    14      netf   =    0
npe    =    2       nps    =   203
ncfn   =    0       ncfl   =    0

Begin backward integration... done.

Final Statistics..

lenrw  =  17316      leniw =   212
llrw   =   8656      lliw  =    80
nst    =    94
nfe    =   117      nfel   =   160
nni    =   113      nli    =   160
nsetups =    20      netf   =    2
npe    =    2       nps    =   253
ncfn   =    0       ncfl   =    0

```





## 4 Parallel tests

The most preeminent advantage of CVODES over existing sensitivity solvers is the possibility of solving very large-scale problems on massively parallel computers. To illustrate this point we present speedup results for the integration and forward sensitivity analysis for an ODE system generated from the following 2-species diurnal kinetics advection-diffusion PDE system in 2 space dimensions. This work was reported in [3]. The PDE is a modification of that described in [4], and takes the form:

$$\frac{dc_i}{dt} = K_h \frac{d^2 c_i}{dx^2} + v \frac{dc_i}{dx} + K_v \frac{d^2 c_i}{dz^2} + R_i(c_1, c_2, t), \quad \text{for } i = 1, 2,$$

where

$$\begin{aligned} R_1(c_1, c_2, t) &= -q_1 c_1 c_3 - q_2 c_1 c_2 + 2q_3(t) c_3 + q_4(t) c_2, \\ R_2(c_1, c_2, t) &= q_1 c_1 c_3 - q_2 c_1 c_2 - q_4(t) c_2, \end{aligned}$$


$K_h$ ,  $K_v$ ,  $v$ ,  $q_1$ ,  $q_2$ , and  $c_3$  are constants, and  $q_3(t)$  and  $q_4(t)$  vary diurnally. The problem is posed on the square  $0 \leq x \leq 20$ ,  $30 \leq z \leq 50$  (all in km), with homogeneous Neumann boundary conditions, and for time  $t$  in  $0 \leq t \leq 86400$  (1 day). The PDE system is treated by central differences on a uniform mesh, except for the advection term, which is treated with a biased 3-point difference formula. The initial profiles are proportional to a simple polynomial in  $x$  and a hyperbolic tangent function in  $z$ .

The solution with CVODES is done with the BDF/GMRES method (i.e. using the CVSPGMR linear solver) and the block-diagonal part of the Newton matrix as a left preconditioner. A copy of the block-diagonal part of the Jacobian is saved and conditionally reused within the preconditioner setup function.

The problem is solved by CVODES using  $P$  processes, treated as a rectangular process grid of size  $p_x \times p_z$ . Each process is assigned a subgrid of size  $n = n_x \times n_z$  of the  $(x, z)$  mesh. Thus the actual mesh size is  $N_x \times N_z = (p_x n_x) \times (p_z n_z)$ , and the ODE system size is  $N = 2N_x N_z$ . Parallel performance tests were performed on ASCI Frost, a 68-node, 16-way SMP system with POWER3 375 MHz processors and 16 GB of memory per node. We present timing results for the integration of only the state equations (column STATES), as well as for the computation of forward sensitivities with respect to the diffusion coefficients  $K_h$  and  $K_v$  using the staggered corrector method without and with error control on the sensitivity variables (columns STG and STG\_FULL, respectively). Run times for a global problem size of  $N = 2N_x N_z = 2 \cdot 1600 \cdot 400 = 1,280,000$  are shown in Fig. 8 and listed below.

$P$	STATES	STG	STG_FULL
4	460.31	1414.53	2208.14
8	211.20	646.59	1064.94
16	97.16	320.78	417.95
32	42.78	137.51	210.84
64	19.50	63.34	83.24
128	13.78	42.71	55.17
256	9.87	31.33	47.95

We note that there was not enough memory to solve the problem (even without carrying sensitivities) using fewer processes.



## References

- [1] A. C. Hindmarsh and R. Serban. Example Programs for CVODE v2.7.0. Technical report, LLNL, 2011. UCRL-SM-208110.
- [2] A. C. Hindmarsh and R. Serban. User Documentation for CVODES v2.6.0. Technical report, LLNL, 2011. UCRL-SM-208111.
- [3] R. Serban and A. C. Hindmarsh. CVODES, the sensitivity-enabled ODE solver in SUNDIALS. In *Proceedings of the 5th International Conference on Multibody Systems, Nonlinear Dynamics and Control*, Long Beach, CA, 2005. ASME.
- [4] M. R. Wittman. Testing of PVODE, a Parallel ODE Solver. Technical Report UCRL-ID-125562, LLNL, August 1996.

